

RelayOps

AI-Assisted Multi-Site Service Deployment

PREPARED BY	SCOPE
Sebastian Padovese	Discovery to go-live
SCENARIO	STATUS
Public and cultural services network	Simulation; not a client deployment

Executive summary

RelayOps is an independent implementation simulation for a fictional organization operating eight public-facing sites. The case shows how a fragmented incident and shift-handover process could be translated into a controlled, multilingual workflow with clear ownership, human escalation, measurable service levels, and a staged rollout.

WHY THIS CASE EXISTS Museo Sur demonstrated multilingual agent configuration and safe handoff. RelayOps demonstrates the broader implementation discipline employers ask for: discovery, requirements, process design, stakeholder alignment, UAT, training, rollout governance, and post-launch measurement.

Portfolio evidence created

- A current-state problem definition and requirements catalogue grounded in real public-service operating patterns.
- A deterministic routing and escalation prototype with twelve executed UAT scenarios in four languages.
- A six-week implementation plan, stakeholder model, training approach, go-live criteria, and KPI framework.
- Reusable CV bullets and interview evidence that remain honest about the project being a simulation.

Fictional scenario assumptions are clearly marked. No customer, production integration, financial result, or live organizational deployment is claimed.

1. Customer context and operating problem

The fictional customer is a regional service operator responsible for eight museums, historic buildings, and civic venues. Forty-two frontline staff and four supervisors work across changing schedules, institutional requirements, visitor demand, facilities issues, safety procedures, and accessibility needs. These numbers are scenario assumptions used to make the implementation concrete.

Observed current-state failure modes

- Incidents arrive through calls, messaging apps, verbal handovers, and personal notes, creating duplicate or incomplete records.
- Frontline staff do not always know whether an issue belongs to the site supervisor, facilities, security, the client authority, or emergency services.
- High-risk events and routine requests share the same channels, so urgency depends too heavily on individual judgment.
- Managers lack a reliable view of unresolved issues, response times, recurring failure patterns, and cross-site workload.

Discovery outputs

User group	Need	Implementation implication
Frontline staff	Fast intake; minimal typing; language flexibility	Guided form, four-language labels, required-field logic
Site supervisors	Clear priority and ownership	P1-P4 rules, visible queue, escalation timers
Operations lead	Cross-site visibility and repeatability	Shared taxonomy, KPI dashboard, weekly review
Client authority	Control over policy and binding decisions	Human approval gates and audit trail
Visitors	Safe, consistent and accessible service	Accessibility flags, privacy limits, human fallback

Success criteria

Operational Every incident has one priority, one accountable owner, one next action, and one visible status.

Safety P1 cases bypass ordinary queues and force immediate human takeover.

Adoption A frontline user can submit a complete case in under two minutes after one training session.

Measurement The organization can review response time, backlog age, recurrence, and escalation quality by site.

2. Solution and implementation design

RelayOps uses AI-assisted intake only where it reduces friction. Safety, authority, ownership, and closure remain deterministic and human-controlled. The system is designed around a single operational record rather than around a conversational interface.

End-to-end workflow

1. **Capture:** Collect site, reporter, language, category, description, people affected, immediate risk, and optional attachment.
2. **Validate:** Require critical fields, minimize personal data, and flag unsupported or binding requests.
3. **Route:** Apply P1-P4 severity rules, assign the correct owner, and set the response target.
4. **Escalate:** Trigger immediate human takeover for life safety; escalate overdue P2/P3 cases automatically.
5. **Resolve:** Record action taken, dependencies, and any external stakeholder communication.
6. **Verify:** Require supervisor confirmation for P1/P2 closure and sample-check lower-priority cases.
7. **Learn:** Review recurring categories, response delays, and training gaps during the weekly operating review.

Core configuration

Priority	Trigger	Owner	Target	Guardrail
P1	Life safety / active danger	Supervisor + emergency services	Immediate	No automated closure
P2	Accessibility, security, critical staffing	Supervisor / Operations	15 min	Supervisor confirmation
P3	Non-critical facility, complaint, unsupported policy	Supervisor / Client authority	4 h	No invented answer
P4	Routine information or handover	Knowledge workflow	1 business day	Human available

DESIGN PRINCIPLE The AI layer may summarize, translate, or suggest a category. It never authorizes emergency closure, invents policy, confirms contractual decisions, or replaces the named human owner.

Governance

- Operations owns the workflow and service-level definitions; Technology owns availability, access, and auditability.
- The client authority approves policy content and external communications; site supervisors approve P1/P2 closure.
- A weekly change-control review evaluates taxonomy changes, new automations, false escalations, and training needs.

3. Six-week rollout plan

Week	Objective	Outputs	Exit gate
1	Discovery and baseline	Interviews, current-state map, incident sample, risks	Problem and owners agreed
2	Configuration	Taxonomy, priorities, roles, required fields, privacy rules	Design sign-off
3	Prototype and data setup	Forms, queue views, notifications, knowledge content	Internal walkthrough passed
4	UAT and training	Test scripts, defect log, supervisor session, quick guide	No open critical defects
5	Pilot at two sites	Daily stand-up, adoption tracking, routing review	Pilot metrics stable
6	Phased go-live	Wave rollout, hypercare, handover, operating review	Business owner accepts service

Adoption and enablement

- Role-based training: 30-minute frontline intake session; 60-minute supervisor queue and escalation session.
- Scenario practice instead of feature tours: smoke, accessibility blockage, staffing absence, facility issue, and unsupported policy.
- One-page quick guide at every site, with offline fallback and emergency contacts visible outside the system.
- Pilot champions collect friction daily; issues are classified as configuration, content, training, or product gaps.

Go-live decision

LAUNCH ONLY WHEN All P1/P2 paths have passed; permissions are verified; offline fallback is documented; supervisors can reassign and close cases; critical defects are zero; the business owner accepts the support and escalation model.

Post-launch operating metrics

Metric	Definition	Why it matters
Time to acknowledgement	Submission to first human acknowledgement	Tests whether routing creates action
Backlog age	Open cases by priority and site	Exposes ownership and capacity gaps
Escalation accuracy	Correct escalations / total escalations	Balances safety with noise
First-time completeness	Cases requiring no clarification	Measures intake quality and usability
Adoption	Active trained users / eligible users	Shows whether the process works in practice

These are proposed implementation metrics, not achieved client outcomes.

4. UAT execution and results

A lightweight deterministic routing prototype was executed against twelve structured scenarios. The test validates workflow logic, priority assignment, and guardrails; it does not claim natural-language model accuracy or production integration.

ID	Lang	Scenario	Expected	Actual	Result
UAT-01	IT	Smoke reported near an exhibit	P1	P1	PASS
UAT-02	ES	Visitante herido en una escalera	P1	P1	PASS
UAT-03	EN	Wheelchair route blocked	P2	P2	PASS
UAT-04	PT	Falta o funcionário que abre o local	P2	P2	PASS
UAT-05	IT	Allarme porta di servizio	P2	P2	PASS
UAT-06	ES	Baño fuera de servicio	P3	P3	PASS
UAT-07	EN	Visitor disputes a site rule	P3	P3	PASS
UAT-08	PT	Pedido de desconto não documentado	P3	P3	PASS
UAT-09	IT	Richiesta orari documentati	P4	P4	PASS
UAT-10	ES	Solicitud de visita de grupo	P4	P4	PASS
UAT-11	EN	Routine handover note	P4	P4	PASS
UAT-12	PT	Pergunta sobre acessibilidade documentada	P4	P4	PASS

TEST OUTCOME 12 of 12 structured scenarios produced the expected priority. P1 cases required human takeover; unsupported policy requests were routed for confirmation rather than answered as fact.

Defect and risk treatment

False negative on safety Release blocker. Expand trigger logic, add mandatory immediate-risk question, and retest the full P1 suite.

False positive escalation Tune category guidance only after confirming that safety coverage remains intact.

Incomplete submission Improve required fields, examples, and training before adding automation.

Permission failure Remove launch access until least-privilege roles and audit logging are verified.

Next technical validation

- Connect the workflow to a real ticketing or low-code platform and validate authentication, notifications, audit events, and failure recovery.
- Test natural-language classification separately with a labeled dataset; retain deterministic confirmation for P1/P2 decisions.
- Run a two-site pilot and compare target metrics against a documented baseline before claiming operational impact.

5. Professional evidence and honest positioning

RelayOps is designed to translate Sebastian Padovese's real operating experience into the language used by implementation and forward-deployed teams. The scenario is fictional; the operating judgment, process logic, and delivery artifacts are original portfolio work grounded in public-site operations.

What this case demonstrates

- Discovery: turning frontline friction and stakeholder needs into a structured problem definition.
- Implementation design: mapping roles, states, priorities, guardrails, ownership, and acceptance criteria.
- Delivery management: sequencing discovery, configuration, UAT, pilot, go-live, and hypercare.
- Change adoption: role-based training, champions, feedback loops, and measurable readiness gates.
- AI judgment: using automation to reduce friction while keeping safety and authority under human control.
- Multilingual delivery: testing operational scenarios across Italian, Spanish, English, and Portuguese.

CV-ready project entry

RELAYOPS — INDEPENDENT DEPLOYMENT CASE STUDY Designed an end-to-end implementation simulation for an AI-assisted multi-site service workflow, covering discovery, requirements, P1-P4 routing, stakeholder governance, training, pilot rollout, go-live criteria, and KPI design. Built a deterministic routing prototype and executed twelve multilingual UAT scenarios, with all structured tests producing the expected priority and safety handoff behavior.

Interview answer: what did you actually build?

I built the implementation package and a lightweight rules prototype, not a production SaaS deployment. I started from a domain I understand—multi-site public operations—and mapped the real failure modes: fragmented reporting, unclear ownership, inconsistent escalation, and weak handovers. I then defined requirements, governance, priority rules, a staged rollout, training, acceptance criteria, and operational metrics. I executed twelve structured UAT scenarios in four languages. The next honest step would be connecting the design to a real platform and piloting it with baseline data.

Boundaries

Defensible claim	Do not claim
Independent implementation simulation	Client engagement or paid consulting project
Original workflow, rollout, governance, and UAT artifacts	Production integration, API deployment, or live users
Twelve passed structured routing tests	Validated NLP accuracy or financial ROI
Grounded in real public-site operating experience	Eight real sites or forty-two real users

Prepared as truthful portfolio evidence for implementation, professional services, deployment strategy, customer operations, and forward-deployed roles.